

Extensible Metatheory Mechanization via Family Polymorphism

With the growing practice of mechanizing language metatheories, it has become ever more pressing that interactive theorem provers make it easy to write reusable, extensible code and proofs. This paper presents a novel language design geared towards extensible metatheory mechanization in a proof assistant. The new design achieves reuse and extensibility via a form of family polymorphism, a seemingly object-oriented idea, that allows code and proofs to be polymorphic to their enclosing families. Our development addresses technical challenges that arise from the underlying language of a proof assistant being simultaneously functional, dependently typed, a logic, and an interactive tool. Our results include (1) a prototypical implementation of the language design as a Coq plugin, (2) a dependent type theory capturing the essence of the language mechanism and its consistency and canonicity results, and (3) case studies showing how the new expressiveness naturally addresses real programming challenges in metatheory mechanization.

1 INTRODUCTION

There is a growing trend among programming languages researchers to use proof assistants to mechanize metatheories. However, the programmer runs into an old problem in the new setting of proof engineering: the Expression Problem (EP) [44].

The EP is a programming challenge that epitomizes the difficulty of writing type-safe, extensible code. To define an expression language that can be reused for future extensions, the programmer faces a fundamental tension [37] between adding new constructors to a data type (e.g., new abstract syntax) and adding new functions over the data type (e.g., new compiler passes).

The EP is well studied in the conventional setting of functional programming and object-oriented (OO) programming. Modern languages, such as Scala [35], have a good supply of linguistic features that offer expressive power to solve the EP.

In contrast, proof assistants offer few linguistic solutions that address the EP. Yet, the challenge of writing extensible, type-safe code is as real, especially for metatheory mechanization. The programmer faces a tension between adding new constructors to an inductive data type (e.g., new abstract syntax) and adding new functions and theorems over the data type (e.g., new meta-theoretical results).

In the Coq proof assistant [9], for instance, inductive types, as well as functions and theorems over inductive types, are closed to extension. So to reuse mechanized metatheories, the common practice is still to copy code and proofs and then modify them in each extension. But having to maintain multiple copies is highly non-modular and antithetical to good software engineering. The programmer could also turn to design patterns and encodings [13, 14, 39, 26, 20]. But they tend to require heavy lifting from the programmer to make code extensible and often lead to non-idiomatic programming styles.


We thus set out to answer the following question: can extensible metatheory mechanization be made easier by having a proof assistant support new *linguistic features* that address the EP?

At the core of many linguistic solutions to the EP is *inheritance*. Inheritance is sometimes interpreted narrowly as a subclass' inheriting methods and instance variables from its superclass. But the language-theoretic essence of inheritance is more general: it is a linguistic approach to incrementally modifying record-like definitions by allowing *late binding* [8].

Language mechanisms including mixins [7], virtual classes [28, 17], virtual types [43], extensible cases [5], etc. are all based on this essential idea of inheritance. In particular, when a language mechanism allows late binding of the meaning of nested types and terms, it is said to support *family polymorphism* [16]: types and terms are polymorphic to a family they are nested within.

Contributions We contribute a language design that integrates family polymorphism into a proof assistant. Because code and proofs are polymorphic to a family they are nested within, they can be inherited and reused by a derived family. Hence, family polymorphism allows for extensible metatheory mechanization.

As an example, the diagram to the right depicts an extensible mechanization of the simply typed λ -calculus (STLC), using family polymorphism. An extension of STLC with products and another with fixpoints can both inherit from the base STLC family: they reuse mechanized metatheories, from abstract syntax all the way to the type-safety theorem, only adding new constructors to inductive types and adding new cases to recursive functions and induction proofs *as needed* by an extension.



Integrating family polymorphism into a dependent type theory for logical reasoning, however, poses significant technical challenges. As we analyze, pillars of dependent type theories—including inductive types, definitional equality, and logical consistency—are all inimical to the kind of extensibility and family polymorphism found in existing OO language designs. Thus, our contributions include novel design recipes for dealing with these challenges and strong metatheoretical guarantees on the underlying logic. Specifically, we make the following contributions.

- We present a language design that enables extensible metatheory mechanization in a higher-order, dependent type theory with inductive types (Section 3). The language design reconciles the expressiveness enabled by family polymorphism with the rigor of a proof assistant, while largely retaining an idiomatic programming style.
- We contribute a prototypical implementation of our language mechanism as a Coq plugin (Section 4). The plugin works by compiling surface-language definitions into Coq modules parameterized by explicit extensibility hooks.
- We capture the essence of the new language mechanism formally by extending Martin–Löf type theory with facilities to express family polymorphism (Section 5). We prove foundational metatheoretical results including consistency and canonicity.
- We present case studies of using our Coq plugin to mechanize language metatheories (Section 6). They show how our language design naturally solves the EP and enables a good amount of reuse and extensibility for mechanizing proofs.

2 DESIGN REQUIREMENTS AND CHALLENGES

Integrating family polymorphism into a proof assistant presents challenges far beyond those found in an object-oriented setting [34, 17, 45], because the underlying programming language is simultaneously functional, dependently typed, a logic, and an interactive tool.

C1. Extensible inductive types vs. exhaustive induction reasoning. Inductive types, generalizing algebraic data types found in functional languages, are a central feature of any proof assistant in use for mechanizing language metatheories. They offer a means to define abstract syntax and inference rules. But unfortunately, inductive types are closed to extension by design.

A family-polymorphism design could potentially support extensible inductive types, by allowing a *derived family* to add new constructors to inductive types inherited from a *base family*. Such a feature would be useful for extending mechanized languages. As an example, Figure 1 shows an STLC extension, the mechanization of which would be made easier by extensible inductive types. Code would be organized into two families, with the derived family inheriting constructors from the base family (e.g., the ellipsis under “Typing rules”) and adding new constructors to model the syntax and semantics of a fixpoint construct.

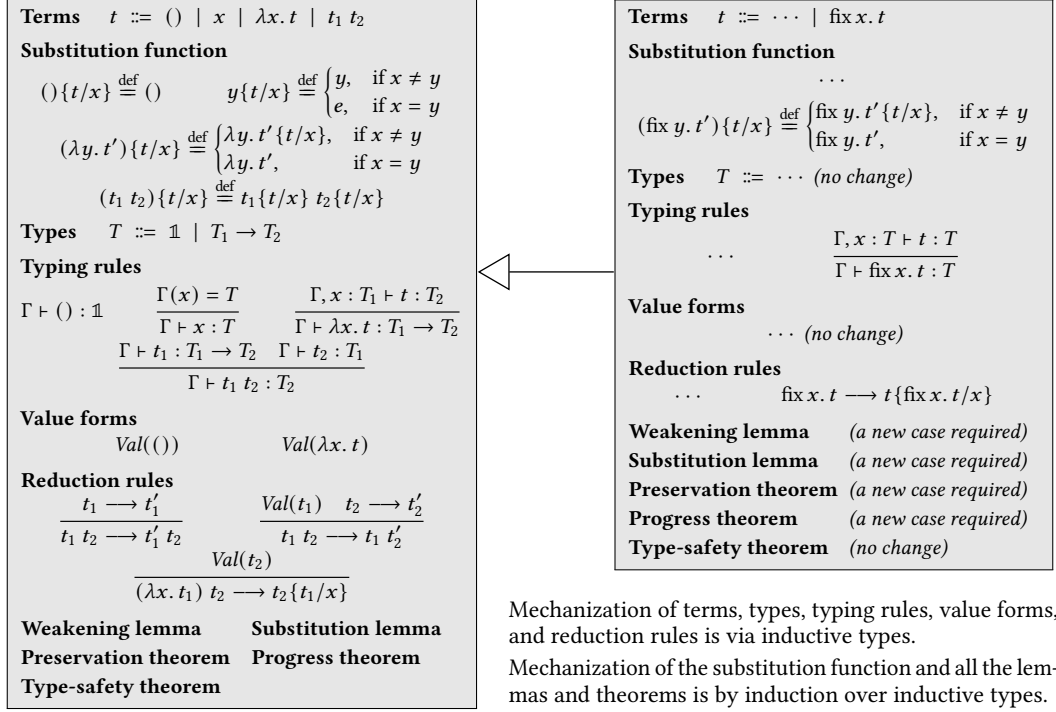


Figure 1. STLC metatheories (left) and its extension with fixpoints (right).

However, there is a tension between extensibility of inductive types and exhaustivity of induction reasoning.¹ In Figure 1, all the lemmas and theorems, as well as the substitution function, require induction (i.e., elimination of inductive types). A language design must enforce that induction remains exhaustive in the face of the new constructors in the derived family. For modularity, the type system should do so without requiring redefinition or rechecking of those cases already handled by the base family.

C2. Late binding vs. definitional equality. Family polymorphism enables modular reuse via late binding: code of a base family can be reused by a derived family, because fields referenced by that code have meanings polymorphic to the enclosing family.

This flexibility, however, prevents *definitional equality* that one takes for granted when programming in a proof assistant. In Figure 1, a proof assistant supporting family polymorphism cannot equate references to the substitution function and its definition (i.e., a pattern match against four cases), as a derived family may modify the definition of the function by adding new cases. Without the ability to unfold the substitution function, how can the programmer even prove the substitution lemma? The problem is compounded by the occasional need in derived families to override fields, which is potentially at odds with being able to use equalities over the fields.

C3. Self reference vs. logical consistency. The language-theoretic essence of late binding is self reference; inheritance and family polymorphism are mechanisms for incrementally modifying

¹The tension reflects a duality between variants and records. For record-like language constructs (e.g., classes and objects), it is safe for an extension to add more fields: existing fields can still be projected. But variant-like constructs (e.g., inductive types) do not automatically enjoy safe, modular addition of constructors: existing pattern matches could potentially become non-exhaustive.

self-referential definitions [8]. However, self reference could easily lead to divergence. Divergence is not a concern for the design of ordinary OO or functional languages, but it would mean logical inconsistency—and hence unsoundness!—for a language aimed at logical reasoning. A family-polymorphism design must tame self reference to guarantee consistency.

C4. User experience and system implementation. Interactive theorem proving and tactic programming are central to a typical programming experience with a proof assistant. A language design integrating family polymorphism should be compatible with these forms of programming. In particular, it should be possible in our system to incrementally navigate through vernacular commands and, moreover, construct proofs with common tactics while getting instant feedback on the proof state, even in the middle of a family definition. Last but not least, in addition to producing theorems, it should be possible for terms defined with families to possess computational content.

3 LANGUAGE DESIGN

We present the key ingredients of our design as an extension to the Coq proof assistant, though we believe the design could be adapted to other proof assistants such as Lean. We call our design and implementation FPOP (family polymorphism for a proof assistant). In this section, we focus on the language design of FPOP. Section 4 describes its implementation as a Coq plugin.

Figure 2 shows how STLC and its extension with fixpoints can be mechanized using FPOP, in a style envisioned in Figure 1. The base family STLC hosts the STLC metatheories, from abstract syntax to the type-safety theorem. Family STLCFix, derived from STLC, makes adjustments as needed by a fixpoints extension: it adds new constructors to the inductive types (FInductive) and adds new cases to the recursive functions (FRecursion) and induction proofs (FInduction). Existing constructors and cases, as well as those definitions and theorems that need no adjustments (ty, env, empty, steps, and typesafe), are automatically inherited and reused. In particular, executing the last command, Check STLCFix. typesafe, displays the type-safety theorem of the fixpoints extension.

3.1 Extensible Inductive Types and Exhaustive Recursion/Induction

Extending inductive types. In family STLCFix, inductive type `tm` further binds the `tm` type in family STLC. It has five constructors: four inherited from STLC and a fifth called `tm_fix`.

Crucially, the meanings of `tm` and its constructors are *late bound*, depending on the family in which they are referenced. Consider `tm_app`. It is defined in family STLC and is thus unaware of `tm_fix`. Yet in family STLCFix, we can use `tm_app` to construct applications of fixpoints, as in `tm_app (tm_fix "f" t1) t2`. This use is justified by `tm_app`'s type, $tm \rightarrow tm \rightarrow tm$. It allows `tm_app` to be applied to anything of type `tm`, which in family STLCFix include those constructed from `tm_fix`.

Ensuring exhaustivity of induction. Ordinarily, an inductive type is not extensible: it is exhaustively generated by its constructors and has no more inhabitants beyond those they construct. This idea is captured by the eliminator (aka *recursor*) associated with an inductive type. For example, if `tm` were defined as an ordinary inductive type, then its eliminator would have the following type:

$$\begin{aligned} \text{tm_rect} : & \forall (P : tm \rightarrow \text{Type}), P \text{ tm_unit} \rightarrow (\forall x, P (tm_var x)) \rightarrow \\ & (\forall x t, P t \rightarrow P (tm_abs x t)) \rightarrow \\ & (\forall t1, P t1 \rightarrow \forall t2, P t2 \rightarrow P (tm_app t1 t2)) \rightarrow \forall t, P t \end{aligned}$$

The eliminator would enable function definitions by recursion and proofs by induction, over `tm`, that exhaustively handle the four cases corresponding to each constructor. The (dependent) return type `P` is called the *motive* of the recursion.

With inductive types made extensible, exhaustivity of induction becomes in jeopardy, however. In particular, recursor `tm_rect` should no longer be allowed, because its type mentions `tm`, whose meaning is late bound, yet `tm_rect` purports to claim $\forall t, P t$ given only four case handlers.

```

197 Family STLC. (* The base STLC *)
198
199 FInductive tm : Set := (* Terms *)
200 | tm_unit : tm
201 | tm_var : id → tm
202 | tm_abs : id → tm → tm
203 | tm_app : tm → tm → tm.
204
205 FRecursion subst on tm (* Substitution function *)
206 motive λ(_ : tm), id → tm → tm.
207 Case tm_unit := λ x t, tm_unit.
208 Case tm_var := λ y x t,
209   if (eqb x y) then t else (tm_var y).
210 Case tm_abs := λ y t' IHT' x t,
211   tm_abs y (if (eqb x y) then t' else IHT' x t).
212 Case tm_app := λ t1 IHT1 t2 IHT2 x t,
213   tm_app (IHT1 x t) (IHT2 x t).
214 End subst.
215
216 FInductive ty : Set := (* Types *)
217 | ty_unit : ty
218 | ty_arrow : ty → ty → ty.
219
220 FDefinition env : Type := id → option ty.
221 FDefinition empty : env := λ _, None.
222
223 FInductive hasty : env → tm → ty → Prop :=
224 | ht_unit : ∀ G, hasty G tm_unit ty_unit
225 | ht_var : ... (* Typing rules *)
226 | ht_abs : ...
227 | ht_app : ....
228
229 FInductive value : tm → Prop := (* Value forms *)
230 | v_unit : value tm_unit
231 | v_abs : ∀ x t, value (tm_abs x t).
232
233 FInductive step : tm → tm → Prop :=
234 | st_app1 : ... (* Reduction rules *)
235 | st_app2 : ...
236 | st_beta : ∀ x t v, value v →
237   step (tm_app (tm_abs x t) v) (subst t x v).
238
239 FDefinition steps := clos_refl_trans step.
240
241 FInduction weakenlem on hasty (* Weaken. lemma *)
242 motive λ G t T (_ : hasty G t T),
243   ∀ G', includedin G G' → hasty G' t T.
244 Case ht_unit. ... Qed. Case ht_var. ... Qed.
245 Case ht_abs. ... Qed. Case ht_app. ... Qed.
246 End weakenlem.
247
248 FInduction substlem on hasty (* Subst. lemma *)
249 motive λ G' t T (_ : hasty G' t T),
250   ∀ G x t' T', G' = extend G x T' →
251   hasty empty t' T' → hasty G (subst t x t') T.
252 Case ht_unit. ... Qed. Case ht_var. ... Qed.
253 Case ht_abs. ... Qed. Case ht_app. ... Qed.
254 End substlem.
255
256 FInduction preserve on hasty (* Preserv. theorem *)
257 motive λ G t T (_ : hasty G t T),
258   G = empty → ∀ t', step t t' → hasty empty t' T.
259 Case ht_unit. ... Qed. Case ht_var. ... Qed.
260 Case ht_abs. ... Qed. Case ht_app. ... Qed.
261 End preserve.
262
263 FInduction progress on hasty (* Progress theorem *)
264 motive λ G t T (_ : hasty G t T),
265   G = empty → value t ∨ ∃ t', step t t'.
266 Case ht_unit. ... Qed. Case ht_var. ... Qed.
267 Case ht_abs. ... Qed. Case ht_app. ... Qed.
268 End progress.
269
270 FTheorem typesafe : (* Type-safety theorem *)
271   ∀ t t' T, steps t t' → hasty empty t T →
272   value t' ∨ ∃ t'', step t' t''.
273 Proof. ... Qed.
274 End STLC.
275
276 Family STLCFix extends STLC.
277   (* STLC extended with fixpoints *)
278
279 FInductive tm : Set +=
280 | tm_fix : id → tm → tm.
281
282 FRecursion subst on tm motive λ _, id → tm → tm.
283 Case tm_fix := λ y t' IHT' x t, ...
284 End subst.
285
286 FInductive hasty : env → tm → ty → Prop +=
287 | ht_fix : ∀ G x t T, hasty (extend G x T) t T →
288   hasty G (tm_fix x t) T.
289
290 FInductive step : tm → tm → Prop +=
291 | st_fix : ∀ x t,
292   step (tm_fix x t) (subst t x (tm_fix x t)).
293
294 FInduction weakenlem on hasty motive ....
295 Case ht_fix. ... Qed.
296 End weakenlem.
297
298 FInduction substlem on hasty motive ....
299 Case ht_fix. ... Qed.
300 End substlem.
301
302 FInduction preserve on hasty motive ....
303 Case ht_fix. ... Qed.
304 End preserve.
305
306 FInduction progress on hasty motive ....
307 Case ht_fix. ... Qed.
308 End progress.
309 End STLCFix.
310
311 Check STLCFix.typesafe.

```

Figure 2. Using FPOP to mechanize STLC and the fixpoints extension, as envisioned in Figure 1.

To reconcile the tension without requiring redefinition or rechecking of case handlers (C1), our design introduces the **FRecursion** and **FInduction** commands. The key idea is to allow case handlers to be added retroactively, should inductive types be extended, and to allow recursion and induction (which are defined in terms of case handlers) to be late bound.

As an example, consider the substitution function `subst`, defined using **FRecursion**. The `on` clause specifies that recursion is over `tm`. The `motive` clause suggests that the recursive function being defined has type $tm \rightarrow id \rightarrow tm \rightarrow tm$. The `subst` function in `STLC` is further bound by the `subst` in `STLCFix`: the four cases from `STLC` are automatically inherited and reused, with `STLCFix` retroactively adding a fifth case to form the new `subst` function. For exhaustivity, it is a static error if the programmer fails to further bind `subst` and define this fifth case. The type system does this check in family `STLCFix` by examining if the inductive type `tm`, over which `subst` is recursively defined, is further bound in the same family.

The **FInduction** command is similar to **FRecursion** but allows cases to be defined in proof mode. Consider `weakenlem` as an example. It is proven by induction on the typing relation `hasty`. Its `motive` clause shows that the lemma reads as $\forall G \ t \ T, \text{hasty } G \ t \ T \rightarrow \forall G', \text{includedin } G \ G' \rightarrow \text{hasty } G' \ t \ T$. Upon entering case `ht_abs`, for instance, Coq enters proof mode with the goal

$$\forall G', \text{includedin } G \ G' \rightarrow \text{hasty } G' \ (tm_abs \ x \ t) \ (ty_arrow \ T1 \ T2)$$

and with the induction hypothesis $\forall G', \text{includedin } (extend \ G \ x \ T1) \ G' \rightarrow \text{hasty } G' \ t \ T2$. The programmer can use tactic programming to discharge the goal. Because the lemma is by **FInduction** on `hasty` and because `STLCFix` adds a constructor `ht_fix` to `hasty`, the programmer is required in family `STLCFix` to extend the proof of `substlem` to handle this extra case.

3.2 Late Binding and Equalities

Late binding of nested names. OO inheritance allows the late binding of method names. Family polymorphism generalizes the power of OO inheritance by allowing the late binding of all names nested within families, including those referring to types. We have seen that late binding of `tm` allows the `tm` constructors in `STLC` to be reused in `STLCFix` to construct terms containing `tm_fix`.

As another example, consider the type of `st_beta` in family `STLC`. It refers to `subst`, whose meaning is late bound. When `st_beta` is inherited into family `STLCFix`, `st_beta` has a type that now refers to the `subst` function in `STLCFix`, where `subst` is defined by handling all the five cases known to that family. Thus, late binding of `subst` allows the derived family to reuse `st_beta` as the β -reduction rule for applications possibly constructed from `tm_fix`.

Importantly, a name is late bound only within a family that defines or further binds it. Outside such families, the name can be accessed only by explicitly specifying a family that contains it. For example, the last line of Figure 2 accesses `typesafe` with a qualifier `STLCFix`. The command prints

```
STLCFix.typesafe:  $\forall t \ t' \ T, \text{STLCFix.steps } t \ t' \rightarrow \text{STLCFix.hasty } \text{STLCFix.empty } t \ T \rightarrow$ 
 $\text{STLCFix.value } t' \ \vee \ \exists t'', \text{STLCFix.step } t' \ t''.$ 
```

In this type, all references to nested names are qualified by `STLCFix`, as desired.

Equality on late bound names. Consider proving the `ht_unit` case of the substitution lemma, `substlem`. The goal is seemingly trivial: the programmer is asked to prove

$$\forall G' \ G \ x \ t' \ T', \ G' = extend \ G \ x \ T' \rightarrow \text{hasty empty } t' \ T' \rightarrow \text{hasty } G \ (\text{subst } tm_unit \ x \ t') \ ty_unit.$$

If `subst` were an ordinary Coq function, then the programmer could discharge the goal with `intros; simpl subst; apply ht_unit`. The Ltac term `simpl subst` unfolds `subst` using its definition and simplifies `subst tm_unit x t'` to `tm_unit`.

But with `subst` being late bound, `subst` cannot and should not be unfolded (C2): the definition of `subst`, as a recursive function, varies across families, yet a derived family should be able to reuse

the proof of the `ht_unit` case even when it has to modify the definition of `subst`. Without the ability to unfold `subst`, how can the programmer make progress in this proof, then?

A key observation is that although late binding prevents *definitional equality* on `subst`, it does not affect *propositional equality*. That is, $\forall x \ t, \text{subst } \text{tm_unit } x \ t = \text{tm_unit}$, as a proposition (`Prop`) about the computational behavior of `subst`, should still hold. After all, how `subst` is defined on `tm_unit` does not vary from a base family to a derived family; what can vary is `subst` itself as a recursive function combining the case handlers.

Based on this insight, FPOP automatically generates a propositional equality for each case handler defined within `FRecursion`, making the equalities and the recursive function available as axioms for use by the rest of the current family. FPOP also provides a tactic `fsimpl` that enables, for instance, simplifying `subst tm_unit x t'` to `tm_unit`. `fsimpl` works by rewriting applications of the axiomatized recursive function using the axiomatized equalities about its computational behaviors.

Note that the definitional equality on `subst` is available *outside* those families that contain `subst`. Within those families, the meaning of `subst` is late bound (i.e., polymorphic to the enclosing family), so only propositional equality is available. In contrast, outside those families, `subst` is always referenced by specifying a family that contains it, as in `STLC.subst` and `STLCFix.subst`. As far as the type checker is concerned, `tm_unit` and `STLCFix.subst tm_unit x t'` are the same thing—the type checker equates them definitionally by unfolding `STLCFix.subst` and performing normalization.

3.3 Overriding

In an OO language, a subclass can override methods of a superclass. Similar expressivity is useful for mechanizing proofs, too. For example, in a derived family, rather than adding new cases to an induction proof, the programmer may prefer overriding the proof entirely, as we observe in our case studies (Section 6).

Overriding is potentially incompatible with having equalities on late bound names, however. Coq distinguishes *opaque* definitions from *transparent* ones. FPOP supports the overriding of opaque definitions, which include most proofs. It is safe to override opaque definitions, because the type checker will never attempt to unfold them. For transparent definitions, the common case is that the programmer does not want to override them. In Figure 2, `env`, `empty`, `subst`, and `steps` are transparent—that is, they are not defined with `Qed`. FPOP treats transparent definitions as non-overridable by default. Thus, the definitional equalities on `env`, `empty`, and `steps`, as well as the propositional equality on `subst`, are available to the type checker for type-checking the families.

FPOP does allow the overriding of transparent definitions explicitly marked as `Overridable` by the programmer. Overriding is made safe by requiring that when an overridable field is overridden, code whose type checking involves unfolding that field should be overridden too. We expect this feature to be used occasionally.

3.4 Sound Logical Reasoning

In Figure 2, just because the `typesafe` theorem in `STLC` is inherited and reused by `STLCFix`, it does not follow that in `STLCFix` the programmer can use `typesafe` to prove progress. If the programmer did, then they would be committing the logical fallacy of circular reasoning: the proof of progress would depend on `typesafe`, yet the proof of `typesafe` depends on progress. Such circularity would easily lead to logical inconsistency. Consider the following two families, where `B` extends `A`:

```
Family A.
  FLemma f : False. Proof. Admitted.
  FLemma g : False. Proof. apply f. Qed.
End A.

Family B extends A.
  FLemma f : False. Proof. apply g. Qed.
End B.
```

`B` overrides lemma `f` by proving it using `g`. Lemma `g` is in turn inherited from family `A`, where it is proven using a late bound reference to lemma `f`. Circularity between `f` and `g` allows proving `False`!

```

344 Family STLCIsorec extends STLC.
345       (* STLC extended with iso-recursive types *)
346 FInductive tm : Set :=
347   | tm_fold : tm → tm | tm_unfold : tm → tm.
348 FRecursion subst on tm motive .... .. End subst.
349 FInductive ty : Set :=
350   | ty_var : id → ty | ty_rec : id → ty → ty.
351 FRecursion tysubst on ty (* Type-level substitution *)
352   motive λ(_ : ty), id → ty → ty.
353 Case ty_unit := .... Case ty_arrow := ....
354 Case ty_var := .... Case ty_rec := ....
355 End tysubst.
356 FInductive hasty : env → tm → ty → Prop :=
357   | ht_fold : ∀ G t α T,
358     hasty G t (tysubst T α (ty_rec α T)) →
359     hasty G (tm_fold t) (ty_rec α T)
360   | ht_unfold : ....
361   ... (* Other adjustments *)
362 End STLCIsorec.
363
364 Family STLCFixIsorec extends STLC
365 using STLCFix, STLCIsorec. (* STLC extended
366                             with fixpoints and iso-recursive types *)
367 End STLCFixIsorec.
368
369 Family STLCProd extends STLC.
370       (* STLC extended with products *)
371 FInductive tm : Set :=
372   | tm_pair : tm → tm → tm
373   | tm_fst : tm → tm | tm_snd : tm → tm.
374 FRecursion subst on tm motive .... .. End subst.
375 FInductive ty : Set :=
376   | ty_prod : ty → ty → ty.
377   ... (* Other adjustments *)
378 End STLCProd.
379
380 Family STLCProdIsorec extends STLC
381 using STLCProd, STLCIsorec. (* STLC extended
382                             with products and iso-recursive types *)
383 FRecursion tysubst on ty
384   motive λ(_ : ty), id → ty → ty.
385 Case ty_prod := .... (* Substitution on product types *)
386 End tysubst.
387 End STLCProdIsorec.
388
389 Family STLCFixProdIsorec extends STLC
390 using STLCFix, STLCProdIsorec. (* STLC extended
391                             with fixpoints, products, and iso-recursive types *)
392 End STLCFixProdIsorec.

```

Figure 3. Composing extensions of STLC.

To ensure the soundness of logical reasoning (C3), the type system requires that in a derived family, the context in which a field is defined be preserved from the base family. In STLC, progress is in the context of `typesafe`. Per the requirement, this relationship must be preserved into `STLCFix`, which prevents the proof of progress from depending on `typesafe` in `STLCFix`.

Note that the requirement still allows a derived family to introduce new declarations into the context of a field. For example, the left column of Figure 3 shows an extension of STLC with iso-recursive types, where `tysubst` is introduced into the context of `hasty`. In the event that `FPOP` cannot infer where the programmer intends to place a new field, annotation is required.

3.5 Composing Families as Mixins

Families can be readily reused to construct larger extensions that *mix in* [7] the functionalities of the individual families. A family like `STLCFix` can be viewed as a family-to-family functor—and hence a mixin, in the sense of [19]—that transforms any family providing the base STLC functionalities (i.e., STLC or a derived family there of) into a new family additionally supporting fixpoints.

In Figure 3, `STLCFixIsorec` is declared as an STLC extension that mixes in `STLCFix` and `STLCIsorec`. The family is declared with minimal verbiage, yet `STLCFixIsorec.typesafe` is automatically a proof of the type-safety theorem of an STLC equipped with fixpoints and iso-recursive types.

Mixin composition is a form of multiple inheritance, which may cause name conflicts in general. `FPOP` requires the programmer to resolve conflicts by overriding conflicted overridable fields.

In the presence of extensible inductive types, mixin composition may also create an obligation to retrofit the mixins with new case handlers. In Figure 3, family `STLCProdIsorec` is composed of two mixins: `STLCProd`, which extends the inductive type `ty` with a new constructor `ty_prod`, and `STLCIsorec`, which introduces a function `tysubst` recursively defined on `ty`. Hence, for exhaustivity, it is required that a composition of `STLCProd` and `STLCIsorec` should additionally handle the `ty_prod` case in `tysubst`.

3.6 Injectivity and Disjointness of Constructors via Partial Recursors

Tactics support for constructors. Coq provides tactics for proving injectivity and disjointness of constructors (i.e., `injection` and `discriminate`). The proof terms generated by the tactics involve exhaustively matching on the constructors of an inductive type, so they do not work for extensible inductive types (C1). In principle, the programmer could use `FInduction` to prove injectivity and disjointness. But this workaround is unsatisfying: it is tedious, it forces the programmer to revisit the induction proofs every time an inductive type is extended, and above all, why should a property like $\neg(\text{tm_var } "x" = \text{tm_abs } "y" \ t)$ have anything to do with `tm_fix`?

To provide a streamlined programming experience (C4), FPOP offers two tactics, `finjection` and `fdiscriminate`. For example, in a proof state that contains a manifestly false assumption $H : \text{tm_var } "x" = \text{tm_abs } "y" \ t$, the programmer can use `fdiscriminate H` to obtain `False` and thus discharge the current goal, just as they would with `discriminate H` if `tm` were not extensible.

Partial recursors. We make the observation that injectivity and disjointness of existing constructors ought to hold regardless of future addition of constructors. This insight motivates the design of *partial recursors*, which power the `finjection` and `fdiscriminate` tactics. Partial recursors are automatically generated for all inductive types defined with `FInductive`.

As analyzed earlier, ordinary recursors, such as `tm_rect`, are impossible within a family in which the name of the inductive type is late bound. However, a key observation is that extensible inductive types still admit a weakened elimination principle where the motive is an `option` type. For example, within family `STLC`, the partial recursor for `tm` has the following type:

```
tm_prect_STLC : ∀ (P : tm → Type), option (P tm_unit) → (∀ x, option (P (tm_var x))) →
  (∀ x t, option (P t) → option (P (tm_abs x t))) →
  (∀ t1, option (P t1) → ∀ t2, option (P t2) → option (P (tm_app t1 t2))) → ∀ t, option (P t)
```

As expected, `tm_prect_STLC` is axiomatized along with four equalities describing its computational behaviors, one for each constructor. For instance, the equality for constructor `tm_abs` is as follows:

```
tm_abs_eq_STLC : ∀ x t P H1 H2 H3 H4,
  tm_prect_STLC P H1 H2 H3 H4 (tm_abs x t) = H3 x t (tm_prect_STLC P H1 H2 H3 H4 t)
```

Importantly, unlike the standard `tm_rect` recursor, the partial recursor `tm_prect_STLC` is compatible with the late binding of `tm` in its type. When `tm_prect_STLC` is inherited into family `STLCFix`, all the previous four equalities still hold, and a trivial, fifth equality is made available:

```
tm_fix_eq_STLC : ∀ x t P H1 H2 H3 H4, tm_prect_STLC P H1 H2 H3 H4 (tm_fix x t) = None
```

Partial recursors appear weaker than ordinary recursors, but there is power in restraint. In particular, they offer a principled, uniform way to derive injectivity and disjointness of constructors, while supporting future extension: they enable injective mappings from a late bound inductive type, like `tm`, to an ordinary inductive type, like \mathbb{N} , the injectivity and disjointness of whose constructors are readily available. For example, in a proof state with the assumption $H : \text{tm_var } "x" = \text{tm_abs } "y" \ t$, running `fdiscriminate H` first applies an injective mapping to both sides of H , obtaining

```
tm_prect_STLC (λ_, IN) (λ_, Some 1) (Some 2) (λ_ --, Some 3) (λ_ --, Some 4) (tm_var "x") =
tm_prect_STLC (λ_, IN) (λ_, Some 1) (Some 2) (λ_ --, Some 3) (λ_ --, Some 4) (tm_abs "y" t)
```

and then rewrites the above equality using the axiomatized computational behaviors of `tm_prect_STLC`, obtaining `Some 1 = Some 3`, from which `False` easily follows. Note that the proof term generated by `fdiscriminate H` in `STLC` is reusable by family `STLCFix`, because the partial recursor and its computational behaviors are compatible with the late binding of `tm`.

Within family `STLCFix`, a second partial recursor (called `tm_prect_STLCFix`) and its computational behaviors are automatically axiomatized, which allows properties of `tm_fix`, such as $\text{tm_fix } x \ t1 = \text{tm_fix } x \ t2 \rightarrow t1 = t2$, to be proved.

4 COMPILING FAMILY POLYMORPHISM TO PARAMETERIZED MODULES

We implement our language design as a Coq plugin. Rather than modifying the Coq kernel, a prototypical implementation as a plugin allows a clearly defined trusted base and improved compatibility with different versions of Coq.

The implementation works by translating programs in FPOP syntax into programs that can be checked and evaluated by Coq. The translation is compatible with interactive theorem proving (C4), in that a family is translated piece by piece, allowing each field to be defined and checked separately. The translation is modular and efficient, in that code compiled for fields of a base family can be shared with derived families without having to be rechecked.

Explicit self parameterization. The spirit of the translation is to take “family polymorphism” literally: every field is translated into a Coq definition that is polymorphic to (i.e., universally quantified over) a representation of its enclosing family. While this universal quantification has been implicit with the FPOP syntax, it has to be made explicit in the translated Coq code.

Figures 4 and 5 illustrate the translation of the STLC and STLCFix families from Figure 2. Fields of a family are translated into parameterized Coq modules (or parameterized module types).

As an example, consider field `env` in family STLC. It is translated into a top-level module named $\text{STLC}^\circ\text{env}$. This module has a parameter called `self` representing the enclosing family: fields of the current family in the context of `env` can be referenced through `self`. In particular, `env` is defined as $\text{id} \rightarrow \text{option } \text{ty}$, where `ty` is a late-bound reference to the `ty` field of the enclosing family. Hence, this reference to `ty` is translated to `self.ty`, which is manifestly polymorphic to the enclosing family. This translation of the `env` field can be shared with a derived family even if it extends `ty` (e.g., STLCProd)—no recompilation is needed because `self.ty` is not tied to any concrete definition of `ty`.

The type of $\text{STLC}^\circ\text{env}$ ’s `self` parameter is $\text{STLC}^\circ\text{env}^\circ\text{Ctx}$, a module type constructed from $\text{STLC}^\circ\text{ty}$ (i.e., the translation of the field before `env`) and its context $\text{STLC}^\circ\text{ty}^\circ\text{Ctx}$. In turn, $\text{STLC}^\circ\text{ty}^\circ\text{Ctx}$ (not shown in Figure 4) is constructed from $\text{STLC}^\circ\text{subst}$, the translation of the field before `ty`, and its context $\text{STLC}^\circ\text{subst}^\circ\text{Ctx}$. Thus, the `self` parameter can be used to reference those and only those fields in the current field’s typing context, which echoes the discussion in Section 3.4.

Translating extensible inductive types. An `FInductive` definition is translated to a parameterized module type. Consider the inductive type `tm`. In Figure 4, it is translated to a top-level module type $\text{STLC}^\circ\text{tm}$ that declares a `tm` type, four functions standing for the constructors, a partial recursor (`tm_prect_STLC`), and the computational behaviors of the partial recursor (e.g., `tm_abs_eq_STLC`).

Importantly, $\text{STLC}^\circ\text{tm}$ merely declares the existence of these names and their types; it does not specify their definitions. Having these names and their types available through the context parameters (`self`) suffices for the translations of the subsequent fields to be type-checked by Coq. Leaving the definitions unspecified enables STLC and STLCFix to instantiate `tm` differently upon `End STLC` and upon `End STLCFix`.

The command `FInductive tm : Set += tm_fix : ...` in family STLCFix is again translated to a module type $\text{STLCFix}^\circ\text{tm}$ (Figure 5). It includes all the names declared by $\text{STLC}^\circ\text{tm}$ via command `Include STLC◦tm(self)`, and additionally declares `tm_fix`, a partial recursor, and related equalities.

Translating recursion and induction. An `FRecursion` definition is translated in two parts: first a module containing the definitions of all the case handlers, and then a module type declaring the existence of the recursive function as well as its computational behaviors.

Consider the translation of `subst` in family STLC. First, a module named $\text{STLC}^\circ\text{subst}^\circ\text{Cases}$ is generated interactively: every time the programmer completes a `Case`, a case handler (e.g., `Def subst◦tm_unit`) is generated and added to the module.

```

491 (* Code emitted upon definition of tm in family STLC *)
492 Module Type STLCotmoCtx.
493 End STLCotmoCtx.
494
495 Module Type STLCotm (self : STLCotmoCtx).
496   Axiom tm : Set.
497   Axiom tm_unit : tm.      Axiom tm_var : id → tm.
498   Axiom tm_abs : id → tm → tm.
499   Axiom tm_app : tm → tm → tm.
500   Axiom tm_prect_STLC : ....
501   Axiom tm_unit_eq_STLC : .... Axiom tm_abs_eq_STLC : ....
502   Axiom tm_abs_eq_STLC : .... Axiom tm_app_eq_STLC : ....
503 End STLCotm.
504
505 (* Code emitted for definition of subst in family STLC *)
506 Module Type STLCosubstoCasesoCtx.
507   Include STLCotmoCtx. Include STLCotm.
508 End STLCosubstoCasesoCtx.
509
510 Module STLCosubstoCases (self : STLCosubstoCasesoCtx).
511   Def substotm_unit := (* emitted upon definition of case *)
512     λ (x : id) (t : self.tm), self.tm_unit.
513   Def substotm_var := ...
514   Def substotm_abs := ... Def substotm_app := ...
515 End STLCosubstoCases.
516
517 Module Type STLCosubstoCtx.
518   Include STLCosubstoCasesoCtx.
519   Include STLCosubstoCases.
520 End STLCosubstoCtx.
521
522 Module Type STLCosubst (self : STLCosubstoCtx).
523   Axiom subst : self.tm → id → self.tm → self.tm.
524   Axiom subst_tm_unit_eq :
525     ∀ x t, self.subst (self.tm_unit) x t =
526     self.substotm_unit x t.
527   Axiom subst_tm_var_eq : ....
528   Axiom ...
529 End STLCosubst.
530
531 (* Code emitted upon definition of ty in family STLC *)
532 ...
533
534 (* Code emitted upon definition of env in family STLC *)
535 Module Type STLCoenvoCtx.
536   Include STLCotyoCtx. Include STLCoty.
537 End STLCoenvoCtx.
538
539 Module STLCoenv (self : STLCoenvoCtx).
540   Def env : Type := id → option self.ty.
541 End STLCoenv.
542
543 (* Code emitted for other fields defined in family STLC *)
544 ...
545
546 (* Code emitted upon conclusion of family STLC *)
547 Module STLC.
548
549   (* Instantiate tm & its constructors *)
550   Inductive tm : Set :=
551   | tm_unit : tm | tm_var : id → tm
552   | tm_abs : id → tm → tm | tm_app : tm → tm → tm.
553   (* Instantiate tm partial recursor & its computational behaviors *)
554   Def tm_prect_STLC := λ P, tm_rect (λ t, option (P t)).
555   Fact tm_unit_eq_STLC : .... reflexivity. Qed.
556   Fact ...
557
558   Include STLCosubstoCases.
559   (* Instantiate subst & its computational behaviors *)
560   Def subst := tm_rect _ substotm_unit
561     substotm_var substotm_abs substotm_app.
562   Fact subst_tm_unit_eq : .... reflexivity. Qed.
563   Fact ...
564
565   (* Instantiate ty, its constructors, partial recursor, etc. *) ...
566   Include STLCoenv. (* Include env *)
567
568   (* Include/Instantiate other fields of family STLC *) ...
569 End STLC.

```

Figure 4. Compilation of family STLC (Figure 2).

Upon `End subst`, a module type named `STLCosubst` is generated. As discussed in Section 3.2, `subst` can be further bound, so its definition is not exposed to the fields that come after it. Accordingly, the translation `STLCosubst` merely declares the types of `subst` and the equalities about its computational behaviors, leaving `subst` undefined and the equalities unproven. The equalities are stated in terms of the case handlers, whose definitions *are* available through the `self` parameter. So Coq can simplify, for example, the type of `subst_tm_unit_eq` to $\forall x t, \text{self.subst } \text{self.tm_unit } x t = \text{self.tm_unit}$. These equalities about the computational behaviors of `subst` will be included and available for use in the translations of the subsequent fields through their `self` parameters.

Importantly, code generated for the case handlers is shared with derived families. In Figure 5, module `STLCFixosubstoCases` reuses—without rechecking—all the case handlers in `STLCosubstoCases` via command `Include STLCosubstoCases(self)`.

The translation of `FInduction` is similar, except that there is no need to register computational behaviors, as `FInduction` proofs are considered opaque.

Translation of further-bindables vs. non-further-bindables. In family `STLC`, field `env` and the case handlers for `subst` are not further-bindable by derived families. In contrast, `tm`, `subst`, and the related equalities can be further bound. The distinction is reflected in the translations. The further-bindable fields are translated to module types that export only types of the fields. The

```

540 (* Code emitted upon definition of tm in STLCFix *)
541 Module Type STLCFix°tm°Ctx.
542 End STLCFix°tm°Ctx.
543 Module Type STLCFix°tm (self : STLCFix°tm°Ctx).
544   Include STLC°tm(self).
545   Axiom tm_fix : id → tm → tm.
546   Axiom tm_fix_eq_STLC : ∀ ..., tm_prect_STLC ... = None.
547   Axiom tm_prect_STLCFix : ....
548   Axiom tm_fix_eq_STLCFix : ....
549 End STLCFix°tm.
550
551 (* Code emitted upon definition of subst in STLCFix *)
552 Module Type STLCFix°subst°Cases°Ctx.
553   Include STLCFix°tm°Ctx. Include STLCFix°tm.
554 End STLCFix°subst°Cases°Ctx.
555
556 Module STLCFix°subst°Cases
557   (self : STLCFix°subst°Cases°Ctx).
558   Include STLC°subst°Cases(self). (* reuse *)
559   Def subst°tm_fix := ... (* translation of new case *)
560 End STLCFix°subst°Cases.
561
562 Module Type STLCFix°subst°Ctx.
563   Include STLCFix°subst°Cases°Ctx.
564   Include STLCFix°subst°Cases.
565 End STLCFix°subst°Ctx.
566
567 Module Type STLCFix°subst (self : STLCFix°subst°Ctx).
568   Include STLC°subst(self).
569   Axiom subst_tm_fix_eq : ....
570 End STLCFix°subst.
571
572 (* Code emitted for other fields defined in STLCFix *) ...
573 (* Code emitted upon conclusion of STLCFix *)
574 Module STLCFix.
575   (* Instantiate tm & its constructors *)
576   Inductive tm : Set :=
577     | tm_unit : tm | tm_var : id → tm
578     | tm_abs : id → tm → tm | tm_app : tm → tm → tm
579     | tm_fix : id → tm → tm.
580   (* Instantiate tm partial recursors & their comp. behaviors *) ...
581   Include STLCFix°subst°Cases.
582   (* Instantiate subst & its computational behaviors *)
583   Def subst := tm_rect _ subst°tm_unit
584     subst°tm_var subst°tm_abs subst°tm_app (* reuse *)
585     subst°tm_fix.
586   Fact subst_tm_unit_eq : .... reflexivity. Qed.
587   Fact ...
588   (* Include ty, its constructors, partial recursor, etc. *) ...
589   Include STLC°env. (* reuse *)
590   (* Include/Instantiate other fields of STLCFix *) ...
591   Include STLC°typesafe. (* reuse *)
592 End STLCFix.
593
594 (* Code emitted upon command Check STLCFix.typesafe *)
595 Check STLCFix.typesafe.

```

Figure 5. Compilation of family STLCFix and the final `Check` command (Figure 2).

non-further-bindable fields are translated to modules that export definitional equalities on the fields. Opaque fields in FPOP can be further bound (Section 3.3); they are translated to Coq modules that export opaque fields.

Eliminating `self` by aggregation. Upon the conclusion of a family definition, a representation of the family is created. For example, module STLC in Figure 4 is generated upon `End STLC`. This module can be viewed as the “fixed point” of the `self`-parameterized translations. The “fixed point” is taken step by step, by adding the translation of each field to this module in the same order as they appear in the family definition.

For the non-further-bindables, the translated modules are directly included (e.g., `Include STLC°env` and `Include STLC°subst°Cases` in Figure 4). The instantiation of `self`s for these modules is implicit, thanks to a Coq nicety: when including a higher-order module, Coq automatically instantiates its parameter with the current interactive module environment. For instance, command `Include STLC°subst°Cases` is successfully executed, because Coq automatically instantiates the `self` parameter using the current module environment, which by construction contains all the fields required by `STLC°subst°Cases°Ctx`.

For the further-bindables, `Axioms` declared in the module types must be instantiated.

- In Figure 4, an inductive type `tm` is generated, instantiating the axiomatized `tm` type and its constructors. The partial recursor `tm_prect_STLC` is defined with the help of `tm_rect`, the recursor Coq generates for `tm`. Computational behaviors of `tm_prect_STLC` are immediate, by `reflexivity`.
- Similarly, `subst` is instantiated by applying the recursor `tm_rect` to the (already included) case handlers. Computational behaviors of `subst` are then immediate, by `reflexivity`.

Module `STLCFix` in Figure 5 is emitted upon `End STLCFix`, in the same way as described above for `STLC`. The translation makes sharing evident. In particular, case handlers compiled for `STLC` are reused to instantiate `subst`, `substlem`, and `alike`. `STLCoenv` and `STLCotypesafe` are also reused in the construction of module `STLCFix`. One may argue that since the first four constructors of `tm` are repeated in `STLCFix`, the translation does not satisfy the modular compilation requirement. We could address this concern by using wrapper types, but we consider restating constructors a small price to pay in return for the clarity and concision of implementation. We emphasize that compiled case handlers, such as `substotm_abs`, are entirely reusable without rechecking, even with restated constructors. Finally, the reference `STLCFix.typesafe` (where `STLCFix` is a family) can simply be translated to `STLCFix.typesafe` (where `STLCFix` is a Coq module), as the last line of Figure 5 shows.

5 FMLTT: A CORE DEPENDENT TYPE THEORY

We contribute FMLTT, a core type theory that extends Martin-Löf dependent type theory (MLTT) [29] with facilities to express family polymorphism while maintaining consistency and canonicity. We acknowledge that the presentation herein is dense for an audience without prior exposure to MLTT, so we summarize the salient points first before delving into the technical details.

Summary. FMLTT is intended as a foundational model. So unlike our programmer-facing plugin, fields automatically axiomatized by the plugin require explicit definitions in FMLTT. FMLTT provides MLTT-style constructs that can be used to express families and family polymorphism. Most notably, it extends MLTT with what we call *linkages*. (Linkages are a namesake of the theoretical device through which prior work [45] models family polymorphism in an OO setting, but the technical details differ significantly from the prior work.)

- Linkages model families; they are like tuples composed of the fields, with field names represented by variable bindings. But there is a twist: linkages support late binding. Unlike dependent tuples where a later component is *existentially* quantified over the earlier ones, a linkage component is *universally* quantified over—and thus polymorphic to—the components preceding it.
- FMLTT also features *linkage transformers* modeling the inductive construction of families from other families. They allow a linkage to be created by inheriting or overriding components of an existing linkage, or by extending an existing linkage with new components.
- Inductive types are modeled as W-types [30] and their extension as overriding.
- Consistency and canonicity of FMLTT are proved by giving semantic interpretations to the syntactic typing judgments.

Brief review of MLTT. Figure 6 presents the syntax and selected typing rules of MLTT (and Figure 7 FMLTT). As is the current trend [10, 2, 21] with MLTT presentations, we use *explicit substitutions* [31, 1]: substitutions γ and their applications (e.g., $T[\gamma]$) are part of the syntax rather than meta-operations. Variables are in the form of *de Bruijn* indices: var_n is the variable bound by the n -th closest enclosing binder. For example, $\lambda x. \lambda y. x$ is $\lambda(\lambda(\text{var}_1))$. Substitutions are typed with the form $\Gamma \vdash \gamma : \Delta$. The idea is that applying γ to terms valid in the context Δ yields terms valid in Γ (**TM/SUB** and **TY/SUB**). The two main forms of substitutions are weakening (**SUB/WK**) and extension (**SUB/EXT**): $t[p^n]$ introduces n free variables into the context of t , and $t[\gamma, t']$ substitutes t' for var_0 in t and then applies γ . For example, rule **TM/SND** states that if t is a dependent pair that has type $\Sigma(A, B)$, then $\text{snd } t$ has type $B[p^0, \text{fst } t]$, where p^0 is the identity substitution (**SUB/ID**).

We use Tarski-style universes [22]: a universe \mathbb{U} is inhabited by the *codes* of types, with $\text{El}(t)$ decoding t (rule **TY/EL**) and $\text{c}(T)$ encoding type T (rule **TM/C**). MLTT has an infinite hierarchy of universes; we omit universe levels in the presentation for conciseness.

A singleton type $\text{S}(t)$ helps expose the definition of a term t in its type (rule **TM/S**) [3, 41].

Contexts	$\Gamma, \Delta ::= \cdot \mid \Gamma, A$
Substitutions	$\gamma ::= \mathbf{p}^n \mid \gamma, t \mid \gamma_1 \circ \gamma_2$
Types	$A, B, T ::= T[\gamma] \mid \mathbb{U} \mid \mathbb{B} \mid \perp \mid \top \mid \Pi(A, B) \mid \Sigma(A, B) \mid \text{Eq}(t_1, t_2) \mid \mathbb{S}(t) \mid \text{El}(t)$
Terms	$t ::= t[\gamma] \mid \text{var}_n \mid \text{c}(T) \mid () \mid \text{tt} \mid \text{ff} \mid \text{if}(t_1, t_2, t_3) \mid \lambda(t) \mid \text{app}(t) \mid$ $(t_1, t_2) \mid \text{fst } t \mid \text{snd } t \mid \text{refl}(t) \mid J(t_1, t_2)$
$\boxed{\Gamma \vdash}$	$\boxed{\Gamma \vdash \gamma : \Delta}$ $\boxed{\Gamma \vdash \gamma_1 \equiv \gamma_2 : \Delta}$ $\boxed{\Gamma \vdash T}$ $\boxed{\Gamma \vdash T_1 \equiv T_2}$ $\boxed{\Gamma \vdash t : T}$ $\boxed{\Gamma \vdash t_1 \equiv t_2 : T}$
(SUB/ID)	(SUB/WK)
$\frac{\Gamma \vdash}{\Gamma \vdash \mathbf{p}^0 : \Gamma}$	$\frac{\Gamma \vdash \mathbf{p}^n : \Delta}{\Gamma, A \vdash \mathbf{p}^{n+1} : \Delta}$
(SUB/EXT)	(TY/SUB)
$\frac{\Gamma \vdash \gamma : \Delta \quad \Delta \vdash A}{\Gamma \vdash \gamma, t : \Delta, A}$	$\frac{\Delta \vdash T}{\Gamma \vdash \gamma : \Delta}$
(TM/SUB)	
$\frac{\Delta \vdash t : T}{\Gamma \vdash t[\gamma] : T[\gamma]}$	
(TMEQ/SUB/ID)	(TM/VAR)
$\frac{\Gamma \vdash t : T}{\Gamma \vdash t[\mathbf{p}^0] \equiv t : T}$	$\frac{\Gamma, A_n, \dots, A_1, A_0 \vdash}{\Gamma, A_n, \dots, A_1, A_0 \vdash \text{var}_n : A_n[\mathbf{p}^{n+1}]}$
(TM/LAM)	(TM/APP)
$\frac{\Gamma \vdash A \quad \Gamma, A \vdash t : B}{\Gamma \vdash \lambda(t) : \Pi(A, B)}$	$\frac{\Gamma \vdash t : \Pi(A, B)}{\Gamma, A \vdash \text{app}(t) : B}$
(TM/PAIR)	(TM/SND)
$\frac{\Gamma \vdash t_1 : A \quad \Gamma \vdash t_2 : B[\mathbf{p}^0, t_1]}{\Gamma \vdash (t_1, t_2) : \Sigma(A, B)}$	$\frac{\Gamma \vdash t : \Sigma(A, B)}{\Gamma \vdash \text{snd } t : B[\mathbf{p}^0, \text{fst } t]}$
(TY/EL)	(TM/C)
$\frac{\Gamma \vdash t : \mathbb{U}}{\Gamma \vdash \text{El}(t)}$	$\frac{\Gamma \vdash T}{\Gamma \vdash \text{c}(T) : \mathbb{U}}$
(TMEQ/C)	(TYEQ/EL)
$\frac{\Gamma \vdash t : \mathbb{U}}{\Gamma \vdash \text{c}(\text{El}(t)) \equiv t : \mathbb{U}}$	$\frac{\Gamma \vdash T}{\Gamma \vdash \text{El}(\text{c}(T)) \equiv T}$
(TY/S)	(TM/S)
$\frac{\Gamma \vdash t : T}{\Gamma \vdash \mathbb{S}(t)}$	$\frac{\Gamma \vdash t : T}{\Gamma \vdash t : \mathbb{S}(t)}$
(TMEQ/S/ETA)	
$\frac{\Gamma \vdash t_1 : T \quad \Gamma \vdash t_2 : \mathbb{S}(t_1)}{\Gamma \vdash t_1 \equiv t_2 : T}$	

Figure 6. Syntax and selected typing rules of MLTT.

Introducing and eliminating inductive types. W-types [30] are a succinct way to model inductive types in MLTT. Together with the identity type $\text{Eq}(\cdot, \cdot)$, they can express a whole host of inductive types [23]. A constructor is given by a pair of types $\Gamma \vdash A$ and $\Gamma, A \vdash B$ (rule **WSIG/ADD**): given two arguments $\Gamma \vdash t_1 : A$ and $\Gamma, B[p^0, t_1] \vdash t_2 : T$, $\text{Wsup}(\tau, t_1, t_2)$ constructs a term of type T (rule **TM/WSUP**). As the rules suggest, our formulation additionally supports *W-type signatures*. A well formed W-type signature $\Gamma \vdash \tau \text{WSig}^n$ is composed of n pairs of types, each of which represents a constructor (rules **WSIG/EMPTY** and **WSIG/ADD**), with $w\pi_1^i(\tau)$ and $w\pi_2^i(\tau)$ giving the i -th pair in τ , and $\text{W}(\tau)$ giving the code of the W-type (**TM/W**).

W-types are eliminated with the form $\text{Wrec}(\tau, \ell, t)$, where t is of a W-type $\text{El}(\text{W}(\tau))$, and ℓ is essentially an n -tuple of case handlers for the n constructors in τ (**TM/WREC**). Each case handler has a type of form $\text{CaseTy}(A, B, T)$, where T is the motive of the recursion (**TYEQ/CASETY**); for simplicity, we model non-dependent motives for recursion. The collection of case handlers ℓ encodes those defined and inherited by an **FRecursion** command in our plugin. We choose to type it with a linkage type $\mathbb{L}(\text{RecSig}(\tau, T))$ to avoid introducing non-dependent n -tuples, which linkages generalize.

Rules **TM/WSUP** and **TM/WREC** require access to τ : the W-type is exhaustively generated by its constructors, and its elimination must exhaustively handle all the constructors in its signature.

In contrast, τ should be hidden from the typing context of any term that does not invoke $\text{Wsup}(\tau, \cdot, \cdot)$ or $\text{Wrec}(\tau, \cdot, \cdot)$, so that the term can be reused—without being rechecked—for a different W-type signature τ' that extends τ with additional constructors. Moreover, the typing of the term should be made parametric to the definitions of those fields that invoke Wsup or Wrec , so that the term can be reused—without being rechecked—when those fields are overridden to support the extended signature τ' . Such abstraction required by family polymorphism is supported via linkages, which we discuss next.

Types	$A, B, T ::= \dots \mid \mathbb{W}(\tau) \mid w\pi_1^i(\tau) \mid w\pi_2^i(\tau) \mid \mathbb{L}(\sigma) \mid \mathbb{P}(\sigma) \mid v\pi_2(\sigma) \mid \text{CaseTy}(A, B, T)$
Terms	$t, s, \ell ::= \dots \mid \mathbb{W}(\tau) \mid \text{Wsup}(\tau, t_1, t_2) \mid \mu^\bullet \mid \mu^+(\ell, t) \mid \text{inh}(h, \ell) \mid \text{Wrec}(\tau, \ell, t) \mid \mu\pi_1(\ell) \mid \mu\pi_2(\ell) \mid v\pi_s(\sigma) \mid \mathbb{P}(\ell) \mid R\pi^i(\ell)$
W-type signatures	$\tau ::= w^\bullet \mid w^+(\tau, A, B) \mid \tau[\gamma] \mid w^-(\tau)$
Linkage signatures	$\sigma ::= v^\bullet \mid v^+(\sigma, s, T) \mid v\pi_1(\sigma) \mid \text{RecSig}(\tau, T) \mid \sigma[\gamma]$
Linkage transformers	$h ::= \text{Identity} \mid \text{Extend}(h, t) \mid \text{Override}(h, t) \mid \text{Inherit}(h) \mid \text{Nest}(h, h')$

$\Gamma \vdash T_1 \equiv T_2$	$\Gamma \vdash t : T$	$\Gamma \vdash t_1 \equiv t_2 : T$	$\Gamma \vdash \tau \text{WSig}^n$	$\Gamma \vdash \sigma \text{LSig}^n$	$\Gamma \vdash h : \sigma_1 \twoheadrightarrow \sigma_2$
<div style="display: flex; justify-content: space-around;"> <div> $\frac{\text{(TM/W)} \quad \Gamma \vdash \tau \text{WSig}^n}{\Gamma \vdash \mathbb{W}(\tau) : \mathbb{U}}$ </div> <div> $\frac{\text{(WSIG/EMPTY)} \quad \Gamma \vdash w^\bullet \text{WSig}^0}{\Gamma \vdash w^\bullet \text{WSig}^0}$ </div> <div> $\frac{\text{(WSIG/ADD)} \quad \Gamma \vdash \tau \text{WSig}^n \quad \Gamma \vdash A \quad \Gamma, A \vdash B}{\Gamma \vdash w^+(\tau, A, B) \text{WSig}^{n+1}}$ </div> <div> $\frac{\text{(TM/WSUP)} \quad \Gamma \vdash \tau \text{WSig}^n \quad \Gamma \vdash t_1 : w\pi_1^i(\tau) \quad \Gamma, w\pi_2^i(\tau)[p^0, t_1] \vdash t_2 : \text{El}(\mathbb{W}(\tau))}{\Gamma \vdash \text{Wsup}(\tau, t_1, t_2) : \text{El}(\mathbb{W}(\tau))}$ </div> </div>					
<div style="display: flex; justify-content: space-around;"> <div> $\frac{\text{(TM/WREC)} \quad \Gamma \vdash \ell : \mathbb{L}(\text{RecSig}(\tau, T)) \quad \Gamma \vdash t : \text{El}(\mathbb{W}(\tau))}{\Gamma \vdash \text{Wrec}(\tau, \ell, t) : T}$ </div> <div> $\frac{\text{(TYEQ/CASETY)} \quad \Gamma \vdash A \quad \Gamma, A \vdash B \quad \Gamma \vdash T}{\Gamma \vdash \text{CaseTy}(A, B, R) \equiv \Pi(A, \Pi(\Pi(B, T[p^2]), T[p^2]))}$ </div> </div>					
<div style="display: flex; justify-content: space-around;"> <div> $\frac{\text{(LSIG/EMPTY)} \quad \Gamma \vdash v^\bullet \text{LSig}^0}{\Gamma \vdash v^\bullet \text{LSig}^0}$ </div> <div> $\frac{\text{(LSIG/ADD)} \quad \Gamma \vdash \sigma \text{LSig}^n \quad \Gamma, A \vdash T}{\Gamma \vdash v^+(\sigma, s, T) \text{LSig}^{n+1}}$ </div> <div> $\frac{\text{(L/EMPTY)} \quad \Gamma \vdash \mu^\bullet : \mathbb{L}(v^\bullet)}{\Gamma \vdash \mu^\bullet : \mathbb{L}(v^\bullet)}$ </div> <div> $\frac{\text{(L/ADD)} \quad \Gamma \vdash \ell : \mathbb{L}(\sigma) \quad \Gamma, A \vdash t : T}{\Gamma \vdash \mu^+(\ell, t) : \mathbb{L}(v^+(\sigma, s, T))}$ </div> <div> $\frac{\text{(TM/INH)} \quad \Gamma \vdash h : \sigma_1 \twoheadrightarrow \sigma_2 \quad \Gamma \vdash \ell : \mathbb{L}(\sigma_1)}{\Gamma \vdash \text{inh}(h, \ell) : \mathbb{L}(\sigma_2)}$ </div> </div>					
<div style="display: flex; justify-content: space-around;"> <div> $\frac{\text{(TYEQ/PK/ADD)} \quad \Gamma \vdash \sigma \text{LSig}^n \quad \Gamma, \mathbb{P}(\sigma) \vdash s : A[p^1]}{\Gamma \vdash \mathbb{P}(v^+(\sigma, s, T)) \equiv \Sigma(\mathbb{P}(\sigma), T[p^1, s])}$ </div> <div> $\frac{\text{(TMEQ/PK/ADD)} \quad \Gamma \vdash \ell : \mathbb{L}(\sigma) \quad \Gamma, \mathbb{P}(\sigma) \vdash s : A[p^1] \quad \Gamma, A \vdash t : T}{\Gamma \vdash \mathbb{P}(\mu^+(\ell, t)) \equiv (\mathbb{P}(\ell), t[p^1, s][p^0, \mathbb{P}(\ell)]) : \mathbb{P}(v^+(\sigma, s, T))}$ </div> </div>					
$\frac{\text{(TMEQ/OV/BETA)} \quad \Gamma \vdash h : \sigma_1 \twoheadrightarrow \sigma_2 \quad \Gamma \vdash \ell : \mathbb{L}(\sigma_1) \quad \Gamma, A_1 \vdash t_1 : T_1 \quad \Gamma, \mathbb{P}(\sigma_1) \vdash s_1 : A_1[p^1] \quad \Gamma, A_2 \vdash t_2 : T_2 \quad \Gamma, \mathbb{P}(\sigma_2) \vdash s_2 : A_2[p^1]}{\Gamma \vdash \text{inh}(\text{Override}(h, t_2), \mu^+(\ell, t_1)) \equiv \mu^+(\text{inh}(h, \ell), t_2) : \mathbb{L}(v^+(\sigma_2, s_2, T_2))}$					

Figure 7. Syntax and selected typing rules of FMLTT.

Family polymorphism via linkages. Family polymorphism requires late binding. In FMLTT, families are expressed through linkages, and late binding of field references is achieved by requiring that typing be polymorphic to the definition of that field.

Well-formedness judgments of *linkage signatures* have form $\Gamma \vdash \sigma \text{LSig}^n$. As rules **LSIG/EMPTY** and **LSIG/ADD** show, a linkage signature is a list of n types. Well-formedness judgments of *linkages* have form $\Gamma \vdash \ell : \mathbb{L}(\sigma)$, where $\mathbb{L}(\sigma)$ is a type formed by σ . As rules **L/EMPTY** and **L/ADD** show, a linkage is a list of n terms, each representing a field of the family modeled by the linkage.

In rule **L/ADD**, the second premise $\Gamma, A \vdash t : T$ is responsible for late binding. Here, A abstracts the context of the current field t , controlling how the types of the fields prior to t are exposed to the typing of t . Crucially, the premise $\Gamma, A \vdash t : T$ makes clear that the typing of t in **L/ADD** is *universally* quantified—rather than *existentially* quantified as is in **TM/PAIR**—over how the fields in t 's context are defined. Late binding enables reuse. A different linkage ℓ' that overrides fields in t 's context (and thus models a derived family) can reuse t —without retyping it—by first projecting t from $\mu^+(\ell, t)$ and then appending it to ℓ' .

The type A in **L/ADD** and **LSIG/ADD**, abstracting the types of the prior fields, does not necessarily contain the same types as those recorded by $\mathbb{L}(\sigma)$, because a field defined as the code $W(\tau)$ of a

736	Family STLC.	$\sigma_0 := \nu^\bullet$	$\ell_0 := \mu^\bullet$	$\cdot \vdash \ell_0 : \mathbb{L}(\sigma_0)$
737				
738	FInductive $\text{tm} :=$	$\sigma_1 := \nu^+(\sigma_0, s_1, T_1)$	$\ell_1 := \mu^+(\ell_0, t_1)$	$\cdot \vdash \ell_1 : \mathbb{L}(\sigma_1) \quad \mathbb{P}(\sigma_0) \vdash s_1 : A_1[p^1]$
739		$\text{self} : [\] \vdash W(\tau_{\text{tm}}) : \mathbb{S}(W(\tau_{\text{tm}}))$		
740	$\text{tm_unit} : \text{tm}$	$\sigma_2 := \nu^+(\sigma_1, s_2, T_2)$	$\ell_2 := \mu^+(\ell_1, t_2)$	$\cdot \vdash \ell_2 : \mathbb{L}(\sigma_2) \quad \mathbb{P}(\sigma_1) \vdash s_2 : A_2[p^1]$
741		$\text{self} : [\text{tm} : \mathbb{S}(W(\tau_{\text{tm}}))] \vdash \text{Wsup}(\tau_{\text{tm}}, \top, \perp) : \text{El}(\text{self}.\text{tm})$		
742		
743	FRecursion subst ...	<i>$W(\tau_{\text{tm}})$ is abstracted into \mathbb{U} in the typing context (i.e., A_6) of the case handler below:</i>		
744				
745	Case $\text{tm_unit} := \dots$	$\sigma_6 := \nu^+(\sigma_5, s_6, T_6)$	$\ell_6 := \mu^+(\ell_5, t_6)$	$\cdot \vdash \ell_6 : \mathbb{L}(\sigma_6) \quad \mathbb{P}(\sigma_5) \vdash s_6 : A_6[p^1]$
746		$\text{self} : [\text{tm} : \mathbb{U}, \text{tm_unit} : \text{El}(\text{tm}), \dots] \vdash \dots : \text{CaseTy}(\top, \perp, \text{El}(\text{self}.\text{tm}) \rightarrow \text{id} \rightarrow \text{El}(\text{self}.\text{tm}))$		
747	Case ... Case ... Case		
748				
749	End subst.	$\sigma_{10} := \nu^+(\sigma_9, s_{10}, T_{10})$	$\ell_{10} := \mu^+(\ell_9, t_{10})$	$\cdot \vdash \ell_{10} : \mathbb{L}(\sigma_{10}) \quad \mathbb{P}(\sigma_9) \vdash s_{10} : A_{10}[p^1]$
750		$\text{self} : [\text{tm} : \mathbb{S}(W(\tau_{\text{tm}})), \text{tm_unit} : \text{El}(\text{tm}), \dots, \text{subst_tm_unit} : \text{CaseTy}(\top, \perp, \dots)] \vdash \lambda t. \text{Wrec}(\tau_{\text{tm}}, \dots, t) : \text{El}(\text{self}.\text{tm}) \rightarrow \text{El}(\text{self}.\text{tm}) \rightarrow \text{id} \rightarrow \text{El}(\text{self}.\text{tm})$		
751		
752	End STLC.	$\cdot \vdash \mathbb{P}(\ell_{53}) : \mathbb{P}(\sigma_{53})$		

Figure 8. FMLTT encoding of the STLC family from Figure 2. Each row $\text{self} : A_i \vdash t_i : T_i$ types a field. The type A_i controls how the typing of field t_i sees the types of the fields prior to t_i .

W-type has to expose different types to different fields that come after it. Later fields that invoke $\text{Wsup}(\tau, \cdot, \cdot)$ or $\text{Wrec}(\tau, \cdot, \cdot)$ should see the concrete signature τ , as rules **TM/WSUP** and **TM/WREC** stipulate. By contrast, τ should be hidden from all other fields, so that they can be reused in a different context where the W-type signature τ is replaced by an extended one τ' .

We use Figure 8 to illustrate. On the left is code excerpted from Figure 2, and on the right is how the corresponding fields are modeled and typed in FMLTT. ℓ_i is the linkage that adds the i -th field t_i with the typing $A_i \vdash t_i : T_i$. For readability, we feel free to use a notation with explicit names in this typing judgment: we name the context as self (cf. Figure 4) and feel free to write its type A_i , which is a dependent tuple type, as a dependent record type with labels.

- The first field t_1 is defined as $W(\tau_{\text{tm}})$ and given the singleton type $\mathbb{S}(W(\tau_{\text{tm}}))$, where τ_{tm} is the W-type signature for tm . This typing is recorded by σ_1 and is thus available in all σ_i 's.
- The next four fields model the four constructors of tm . Constructor tm_unit is modeled as $\text{Wsup}(\tau_{\text{tm}}, \top, \perp)$ and typed with $\text{El}(\text{self}.\text{tm})$, where self stands for the typing context containing the first field tm . The W-type signature τ_{tm} is exposed in this typing context; that is, tm has type $\mathbb{S}(W(\tau_{\text{tm}}))$. So $\text{El}(\text{self}.\text{tm})$ and $\text{El}(W(\tau_{\text{tm}}))$ can be equated, as required by rule **TM/WSUP**.
- Similarly, τ_{tm} is exposed in the typing context of t_{10} , which models the recursive function subst by invoking the recursor $\text{Wrec}(\tau_{\text{tm}}, \cdot, \cdot)$. Like subst , partial recursors in FPOP are axiomatized by the plugin (Sections 3.2 and 3.6), and they can similarly be defined in FMLTT using Wrec .
- By contrast, τ_{tm} is hidden from the typing of all other fields. Their typing should depend on the knowledge that tm has type \mathbb{U} , rather than $\mathbb{S}(W(\tau_{\text{tm}}))$, so that they can be reused in a context where tm is defined as $W(\tau'_{\text{tm}})$, where τ'_{tm} extends τ_{tm} with additional constructors. For example, in Figure 8, the typing of the case handlers of subst (e.g., t_6) is oblivious to the definition of tm —it sees only $\text{tm} : \mathbb{U}$ —so the case handlers can be reused by a linkage modeling **STLCfix**.

The third premise $\Gamma, \mathbb{P}(\sigma) \vdash s : A[p^1]$ in rules **L/ADD** and **LSIG/ADD** is responsible for hiding W-type signatures. Here, $\mathbb{P}(\sigma)$ is a dependent tuple type (rule **TYEQ/PK/ADD**) that packages the types of all fields preceding the current field t , and s is the term turning the tuple $\mathbb{P}(\ell)$ of type $\mathbb{P}(\sigma)$ (rule **TMEQ/PK/ADD**) into a new tuple of type A that hides W-type signatures behind \mathbb{U} , if necessary.

It is straightforward to find the s that fits the bill, though this process is not automated in FMLTT. In particular, when no hiding is needed, s is var_0 and A is $\mathbb{P}(\sigma)$. In Figure 8, s_1 , s_2 , and s_{10} are var_0 , and s_6 is from $[\text{tm} : \mathbb{S}(W(\tau_{\text{tm}})), \text{tm_unit} : \text{El}(\text{tm}), \dots]$ to $[\text{tm} : \mathbb{U}, \text{tm_unit} : \text{El}(\text{tm}), \dots]$.

When a family is concluded (e.g., **End** STLC), a linkage ℓ containing all the fields is available (e.g., ℓ_{53} in Figure 8). Fields of the family can then be accessed by projecting them out of the tuple $\mathbb{P}(\ell)$.

Linkage transformers. Inheritance and code reuse can already be expressed through the projection of fields out of linkages and their inclusion into new linkages. To make common patterns of linkage manipulations more convenient, FMLTT provides a “library” of *linkage transformers*, whose well-formedness judgments have form $\Gamma \vdash h : \sigma_1 \twoheadrightarrow \sigma_2$. The idea is that applying h to a linkage of type $\mathbb{L}(\sigma_1)$ yields a linkage of type $\mathbb{L}(\sigma_2)$ (**TM/INH**).

Derived families can be modeled as linkage transformers inductively constructed from the introduction forms **Identity**, **Extend**(h, t), **Override**(h, t), **Inherit**(h), etc. Figure 7 shows the β -rule of an exemplary transformer, **TMEQ/OV/BETA**. It states that applying the transformer **Override**(h, t_2) to a linkage of form $\mu^+(\ell, t_1)$ overrides the linkage’s last field t_1 with t_2 . For instance, the construction below shows that **Override**(**Identity**, $W(\tau'_{\text{tm}})$) is used as the first step in creating a linkage transformer modeling a derived family that overrides τ_{tm} with an extended signature τ'_{tm} :

```
Family STLCFix extends STLC.   $h_0 := \text{Identity} \quad \cdot \vdash h_0 : v^\bullet \twoheadrightarrow v^\bullet$ 
FInductive tm += ...           $h_1 := \text{Override}(h_0, W(\tau'_{\text{tm}})) \quad \cdot \vdash h_1 : v^+(\nu^\bullet, s_1, \mathbb{S}(W(\tau_{\text{tm}}))) \twoheadrightarrow v^+(\nu^\bullet, s'_1, \mathbb{S}(W(\tau'_{\text{tm}})))$ 
```

A supplemental appendix sketches how the other introduction forms of linkage transformers can be used to model the construction of a derived family as a linkage transformer.

The complete formalization. The definitive version containing all the rules in FMLTT is stated in terms of a metalanguage supporting quotient inductive-inductive types (QIITs) [33, 2]. QIITs support equality constructors, which facilitate the expression of conversion rules (e.g., **TMEQ/SUB/ID**). The formalization is available as a supplemental file in Agda syntax, though it is not checked by Agda because Agda does not support QIITs natively. An effort to mechanize the formalization is underway, using Agda’s **REWRITE** pragma to replace the limited uses of equality reflection. We also typeset a more complete set of rules than the space allows and present it in a supplemental appendix.

Consistency. One of the most fundamental properties of a dependent type theory is consistency.

THEOREM 5.1 (CONSISTENCY). *The typing judgment $\cdot \vdash t : \perp$ is not derivable for any term t .*

Consistency says that the type \perp is not inhabited. Thus, it is safe to use the type theory for logical reasoning, as not every proposition is trivially provable.

Theorem 5.1 is a consequence of a metacircular interpretation of all the well-formedness rules, which, in particular, interprets \perp in FMLTT as an uninhabited type in the metalanguage. The construction of this interpretation follows the standard model of Altenkirch and Kaposi [2], extending it to handle linkages. A linkage $\mu^+(\ell, t)$ can be interpreted into a non-dependent pair where the second component is a dependent function (modeling late binding), as rule **L/ADD** indicates. We refer the reader to a supplemental file in Agda syntax for the complete proof.

Canonicity. A second property we prove is canonicity.

THEOREM 5.2 (CANONICITY). *If $\cdot \vdash t : \mathbb{B}$, then either $\cdot \vdash t \equiv \text{tt} : \mathbb{B}$ or $\cdot \vdash t \equiv \text{ff} : \mathbb{B}$.*

This canonicity theorem says that every closed term of the ground type \mathbb{B} is convertible to one of the canonical forms **tt** and **ff**. When the canonicity theorem is proved in a constructive metalogic, its proof amounts to a normalization function for closed terms of the ground type. So canonicity serves to justify the computational nature of the type theory.

We prove [Theorem 5.2](#) by constructing a logical-relations model for the well-formedness rules, following prior approaches [10, 40, 25]. In particular, a closed, well formed type $\vdash T$ is interpreted into a logical predicate on closed terms: the predicate includes all closed, “reducible” terms of type T . The proof is available as a supplemental file in Agda syntax.

6 CASE STUDIES

Type safety of STLCs. The first case study is the mechanization of the type safety theorem of STLC and those of its extensions, which has been occurring in the examples in this paper. The code base is ported from Software Foundations [36]. The base STLC family consists of about 400 LOC. Lines of code in each of the four derived families (Y , \times , $+$, and μ in the Venn diagram) vary from 100 to 250, largely depending on how many constructors they add to the inductive types. The linguistic nature of our approach allows us to retain a programming style similar to the original proofs in Software Foundations.

Using individual families to organize the mechanization of individual language features leads to a modular design that also facilitates code reuse. Individually developed features can be easily composed (as mixins) to form new STLC variants.

Composing features can lead to *feature interactions* [4]: features working correctly in isolation may require coordination when composed. For example, composing `STLCIsorec` and `STLCProd` (Figure 3) creates a need to extend `tysubst` to handle `ty_prod`, which the type-checker enforces.


Elimination of inductive types defined via `FInductive` is mostly via the `FRecursion` and `FInduction` commands. An exception is a handful of trivial “inversion lemmas”. For example, consider the lemma $\forall t, \neg \text{step } \text{tm_true } t$ stating that `tm_true` is irreducible. If `step` were an ordinary inductive type, then it could be proved in Coq simply by `intros t H; inversion H`. But `step` is extensible. So one way to prove the lemma is by `FInduction` on `step` and verifying that a derived family does not accidentally make `tm_true` reducible. We observe that it is lighter-weight to use overriding (Section 3.3) instead: the programmer can specify that the proof of the lemma should be overridden in any derived family that further binds `step`, and in return, they are permitted to treat `step` as an ordinary inductive type in the proof and thus use `inversion` to prove it. The plugin then automatically tries the same proof script in a derived family to override the proof. Although proof scripts rather than proof terms are reused, this practice seems justified by the triviality of the lemmas and the terseness of the proof scripts.

The current plugin implementation does not yet support mutually inductive types or mutual induction, which are useful for modeling languages much more complex than STLCs. However, we believe accommodating mutually inductive types does not pose theoretical challenges.

Abstract interpreters for imperative languages. Our second case study is a mechanization of abstract interpreters for simple imperative languages. In addition to the metatheories, this case study produces abstract interpreters that are directly ready for program extraction.

The code is organized into four families. A base family, `Imp`, defines via `FInductive` the abstract syntax of a while-language with pure expressions and impure statements. The semantics is given by an interpreter defined as a CEK-style abstract machine [18] and parameterized by a fuel value. Family `Imp` defines the interpreter via `FRecursion`. It contains ~200 LOC.

A second family, `ImpGAI`, extends `Imp` and uses ~550 LOC. It exports a generic framework for deriving abstract interpreters with partial-correctness guarantees. Soundness of the abstract interpreter, `analyze`, is stated with respect to the interpreter, `eval`, inherited from `Imp`. The theorem says



that the abstraction relation $RState$ over a concrete state S and an abstract state $absS$ is preserved by the analysis:

$$\forall \text{ stmt fuel } S \text{ absS}, RState \text{ } S \text{ absS} \rightarrow RState (\text{eval fuel stmt } S) (\text{analyze fuel stmt absS})$$


analyze is defined via **FRecursion**, and the soundness theorem is proved via **FInduction**. This family leaves fields representing the abstract domain, the abstraction relation, monotonicity of transfer functions, etc. largely unspecified or unproven—a derived family can further bind these “parameters” by overriding appropriate fields (and also possibly extend the abstract syntax), to create a sound, runnable abstract interpreter for a (possibly extended) while-language.

The other two families then extend **ImpGAI**. Family **ImpTI** (~200 LOC) is an abstract interpreter that does type inference [11]. Family **ImpCP** (~300 LOC) extends the abstract syntax with natural number arithmetic, and further binds the generic abstract interpreter to perform constant propagation.

Our implementation of family polymorphism retains the computation ability of the host proof assistant; the two abstract interpreters in **ImpTI** and **ImpCP** are ready for program extraction to OCaml. We put them to test over queries.

We note that our plugin cannot yet allow the expression of two derived abstract interpreters nested within the same family. Work on nested inheritance [34, 45] points to a direction to further increase the expressive power of our language design.

In addition to the two case studies above, we also use extensible inductive types for modeling extensible context-free grammars and derive decision procedures for them.



7 RELATED WORK

Approaches to modular mechanization or proof reuse exist, with different focuses and trade-offs.

Encodings based on product lines or DTC. Delaware et al. [13] engineer product lines of theorems and proofs built from feature modules. Feature composition is manual, which seems to have motivated later approaches based on DTC.

The design pattern *data types à la carte* (DTC) introduced for Haskell by Swierstra [42] is a source of inspiration for many approaches [14, 15, 39, 26, 20]. The original DTC encoding requires type-level general recursion that fails the strict positivity check imposed by proof assistants including Coq and Agda. Delaware et al. [14] introduce *meta-theory à la carte* (MTC): it overcomes the problem by using Church encodings for data types and using Mendler-style folds for evaluation, though it requires **Set** impredicativity. Feature composition is automated through heavy use of type classes. The framework is implemented as a Coq library. Schwaab and Siek [39] adapt DTC to Agda by considering a restricted class of functors that admit least fixed points. Keuchel and Schrijvers [26] use datatype-generic programming techniques for the underlying representation of type-level fixed points and avoid **Set** impredicativity.

All of these approaches are largely extralinguistic, in that they work within the confine of the language offered by a proof assistant, which comes with trade-offs. On the one hand, they can be conveniently distributed as libraries, and the encoding can be more easily adapted for new purposes. For example, MTC [14] has been applied to implementing composable program adverbs [27] and been adapted to support unanticipated feature extensions that require type changes [15].

On the other hand, the extralinguistic nature of the approaches tend to lead to non-idiomatic code and offset their user-friendliness. In particular, because data types have to be encoded (rather than expressed through natively supported inductive types), the resulting code is less readable and the resulting programming style inaccessible to non-experts. In addition, extra programmer effort may be required, such as having to manually prove additional well-formedness conditions.

Forster and Stark [20] introduce *Coq à la carte*. It still follows DTC, but rather than embracing DTC’s use of generic fixed points, it considers specific instantiations instead. The resulting mechanism appears more streamlined than prior *à la carte* approaches particularly for its extensive tool support for generating boilerplate code. But even with the tool support, components (e.g., `subst1lem`) of individually developed feature extensions have to be composed separately by invoking the tool.

Our approach addresses the expression problem by extending the linguistic facilities offered by a proof assistant. Families, in particular, offer an organizational advantage. They allow grouping and coevolving related types, functions, and proofs without explicit parameterization; all further-bindable fields are automatically extensibility hooks. Because family polymorphism does not require explicit parameterization or complex encodings, the resulting programming experience and code are accessible to the working Coq programmer. The more OO aspects of family polymorphism, such as the ability to use families as mixins and the ability to grow a series of mechanized languages in integral increments, also facilitate extensibility and reuse with minimal programmer effort.

Proof reuse and proof repair. Boite [6] addresses proof reuse specifically in response to inductive types extended with new constructors. Proof reuse is via a tactic that adapts the original proof to the extended inductive type while generating proof obligations, so rechecking of proof terms is entailed. The design requires distinct names for a base inductive type and its extensions (including distinct names for constructors), while family polymorphism in our approach allows names to be late bound.

Mulhern [32] introduces a heuristic approach that allows proofs for multiple small languages to be combined to yield proofs for composite languages, as long as the proof structure follows the same pattern. Johnsen and Lüth [24] enable proof reuse in Isabelle by adapting theorems from one setting for reuse in another: proof terms are transformed by first explicitly stating all assumptions, and then abstracting over function symbols and type constants.

Pumpkin Pi [38] is a Coq plug-in that helps repair proofs broken by changes in type definitions. Its decompiler from proof terms to proof scripts prioritizes suggesting useful tactics over soundness. While Pumpkin Pi focuses on refactoring existing proofs in response to changed definitions, our solution can be viewed as an effort to preempt refactoring by enabling the programmer to write code that has built-in hooks for future extensions.

Solutions to the expression problem abound. Almost all involve some form of either explicit or implicit parameterization as extensibility hooks. Our approach is the first that applies family polymorphism [16], an idea from the context of OO languages, to the context of mechanized proofs.

Blume et al. [5] address the expression problem for a core subset of Standard ML by combining explicitly coded open recursion with a mechanism that allows pattern-matching cases to be defined separately and combined later. Our `FRecursion` and `FInduction` commands achieve a similar functionality, with families making open recursion implicit and bestowing organizational power.

8 CONCLUSION

It is hard to write modular, extensible code and proofs. We have presented a solution that equips a proof assistant with linguistic facilities for family polymorphism. The language design ensures that the expressive power brought by family polymorphism is in harmony with the strictness of a proof assistant, while incurring low cognitive overhead and allowing an idiomatic programming experience. We implement the design via a translation to Coq and demonstrate its applicability using case studies. A novel dependent type theory formalizes the essence of the language mechanism and is shown to enjoy consistency and canonicity. Generalizing the mechanism to support nested family polymorphism will make extensibility scale to sizable proof engineering efforts and thus make exciting directions for future work.

REFERENCES

- [1] Martín Abadi, Luca Cardelli, Pierre-Louis Curien, and Jean-Jacques Lévy. 1989. Explicit substitutions. In *ACM SIGPLAN Symp. on Principles of Programming Languages (POPL)*. <https://doi.org/10.1145/96709.96712>
- [2] Thorsten Altenkirch and Ambrus Kaposi. 2016. Type theory in type theory using quotient inductive types. In *ACM SIGPLAN Symp. on Principles of Programming Languages (POPL)*. <https://doi.org/10.1145/2837614.2837638>
- [3] David Aspinall. 1995. Subtyping with singleton types. In *Int'l Workshop on Computer Science Logic*. <https://doi.org/10.1007/BFb0022243>
- [4] Don Batory, Peter Höfner, and Jongwook Kim. 2011. Feature interactions, products, and composition. In *ACM Int'l Conf. on Generative Programming and Component Engineering (GPCE)*. <https://doi.org/10.1145/2047862.2047867>
- [5] Matthias Blume, Umut A. Acar, and Wonseok Chae. 2006. Extensible programming with first-class cases. In *ACM SIGPLAN Conf. on Functional Programming (ICFP)*. <https://doi.org/10.1145/1159803.1159836>
- [6] Olivier Boite. 2004. Proof reuse with extended inductive types. In *Int'l Conf. on Theorem Proving in Higher Order Logics*.
- [7] Gilad Bracha and William Cook. 1990. Mixin-based inheritance. In *ACM SIGPLAN Conf. on Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA)*. <https://doi.org/10.1145/97945.97982>
- [8] William R. Cook, Walter L. Hill, and Peter S. Canning. 1990. Inheritance is not subtyping. In *ACM SIGPLAN Symp. on Principles of Programming Languages (POPL)*.
- [9] Coq [n.d.]. The Coq proof assistant. <https://coq.inria.fr>.
- [10] Thierry Coquand. 2019. Canonicity and normalization for dependent type theory. *Theoretical Computer Science* 777, C (2019). <https://doi.org/10.1016/j.tcs.2019.01.015>
- [11] Patrick Cousot. 1997. Types as abstract interpretations. In *ACM SIGPLAN Symp. on Principles of Programming Languages (POPL)*. <https://doi.org/10.1145/263699.263744>
- [12] Nicolaas Govert de Bruijn. 1972. Lambda calculus notation with nameless dummies, a tool for automatic formula manipulation, with application to the Church-Rosser theorem. *Indagationes Mathematicae* 75, 5 (1972). [https://doi.org/10.1016/1385-7258\(72\)90034-0](https://doi.org/10.1016/1385-7258(72)90034-0)
- [13] Benjamin Delaware, William Cook, and Don Batory. 2011. Product lines of theorems. *ACM SIGPLAN Conf. on Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA)*. <https://doi.org/10.1145/2076021.2048113>
- [14] Benjamin Delaware, Bruno C. d. S. Oliveira, and Tom Schrijvers. 2013. Meta-theory à la carte. In *ACM SIGPLAN Symp. on Principles of Programming Languages (POPL)*. <https://doi.org/10.1145/2429069.2429094>
- [15] Benjamin Delaware, Steven Keuchel, Tom Schrijvers, and Bruno C.d.S. Oliveira. 2013. Modular monadic meta-theory. In *ACM SIGPLAN Conf. on Functional Programming (ICFP)*. <https://doi.org/10.1145/2500365.2500587>
- [16] Erik Ernst. 2001. Family polymorphism. In *European Conf. on Object-Oriented Programming (ECOOP)*.
- [17] Erik Ernst, Klaus Ostermann, and William R. Cook. 2006. A virtual class calculus. In *ACM SIGPLAN Symp. on Principles of Programming Languages (POPL)*. <https://doi.org/10.1145/1111037.1111062>
- [18] Matthias Felleisen and Daniel P. Friedman. 1986. *Control Operators, the SECD-machine, and the λ -calculus*. Technical Report. Computer Science Department, Indiana University.
- [19] Matthew Flatt, Shriram Krishnamurthi, and Matthias Felleisen. 1998. Classes and mixins. In *ACM SIGPLAN Symp. on Principles of Programming Languages (POPL)*. <https://doi.org/10.1145/268946.268961>
- [20] Yannick Forster and Kathrin Stark. 2020. Coq à la carte: A practical approach to modular syntax with binders. In *ACM SIGPLAN Conf. on Certified Programs and Proofs (CPP)*. <https://doi.org/10.1145/3372885.3373817>
- [21] Daniel Gratzer, Jonathan Sterling, and Lars Birkedal. 2019. Implementing a modal dependent type theory. *Proc. of the ACM on Programming Languages (PACMPL)* 3, ICFP (2019).
- [22] Martin Hofmann. 1997. Syntax and semantics of dependent types. In *Extensional Constructs in Intensional Type Theory*.
- [23] Jasper Hugunin. 2020. Why not W?. In *Int'l Conf. on Types for Proofs and Programs (TYPES)*. <https://doi.org/10.4230/LIPIcs.TYPES.2020.8>
- [24] Einar Broch Johnsen and Christoph Lüth. 2004. Theorem reuse by proof term transformation. In *Int'l Conf. on Theorem Proving in Higher Order Logics*. https://doi.org/10.1007/978-3-540-30142-4_12

- [25] Ambrus Kaposi, Simon Huber, and Christian Sattler. 2019. Gluing for type theory. In *Int'l Conf. on Formal Structures for Computation and Deduction (FSCD)*.
- [26] Steven Keuchel and Tom Schrijvers. 2013. Generic datatypes à la carte. In *9th ACM SIGPLAN Workshop on Generic Programming*.
- [27] Yao Li and Stephanie Weirich. 2022. Program adverbs and Tlön embeddings. *Proc. of the ACM on Programming Languages (PACMPL)* 6, ICFP (2022). <https://doi.org/10.1145/3547632> arXiv:2207.05227
- [28] Ole L. Madsen and Birger Moller-Pedersen. 1989. Virtual classes: A powerful mechanism in object-oriented programming. In *ACM SIGPLAN Conf. on Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA)*. <https://doi.org/10.1145/74877.74919>
- [29] Per Martin-Löf. 1982. Constructive mathematics and computer programming. In *Studies in Logic and the Foundations of Mathematics*. Vol. 104.
- [30] Per Martin-Löf. 1984. *Intuitionistic Type Theory: Notes by Giovanni Sambin of a Series of Lectures Given in Padua, June 1980 (Studies in Proof Theory)*.
- [31] Per Martin-Löf. 1992. Substitution calculus. Notes from a lecture given in Göteborg.
- [32] Anne Mulhern. 2006. Proof weaving. In *1st Informal ACM SIGPLAN Workshop on Mechanizing Metatheory*.
- [33] Fredrik Nordvall Forsberg and Anton Setzer. 2010. Inductive-inductive definitions. In *Int'l Workshop on Computer Science Logic*.
- [34] Nathaniel Nystrom, Stephen Chong, and Andrew C. Myers. 2004. Scalable extensibility via nested inheritance. In *ACM SIGPLAN Conf. on Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA)*. <https://doi.org/10.1145/1028976.1028986>
- [35] Martin Odersky and Matthias Zenger. 2005. Scalable component abstractions. In *ACM SIGPLAN Conf. on Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA)*. <https://doi.org/10.1145/1094811.1094815>
- [36] Benjamin C. Pierce, Arthur Azevedo de Amorim, Chris Casinghino, Marco Gaboardi, Michael Greenberg, Cătălin Hrițcu, Vilhelm Sjöberg, and Brent Yorgey. 2022. Software foundations: Volume 2 (programming language foundations). (2022). Version 6.2.
- [37] John C. Reynolds. 1975. User-defined types and procedural data structures as complementary approaches to data abstraction. In *New Directions in Algorithmic Languages*.
- [38] Talia Ringer, RanDair Porter, Nathaniel Yazdani, John Leo, and Dan Grossman. 2021. Proof repair across type equivalences. In *ACM SIGPLAN Conf. on Programming Language Design and Implementation (PLDI)*. <https://doi.org/10.1145/3453483.3454033>
- [39] Christopher Schwaab and Jeremy G. Siek. 2013. Modular type-safety proofs in Agda. In *7th Workshop on Programming Languages Meets Program Verification*. <https://doi.org/10.1145/2428116.2428120>
- [40] Jonathan Sterling. 2019. Algebraic type theory and universe hierarchies. (2019). arXiv:1902.08848
- [41] Christopher Allan Stone. 2000. *Singleton kinds and singleton types*. Ph.D. Dissertation. Carnegie Mellon University.
- [42] Wouter Swierstra. 2008. Data types à la carte. *Journal of Functional Programming (JFP)* 18, 4 (2008). <https://doi.org/10.1017/S0956796808006758>
- [43] Kresten Krab Thorup. 1997. Genericity in Java with virtual types. In *European Conf. on Object-Oriented Programming (ECOOP)*.
- [44] Philip Wadler et al. 1998. The expression problem. <http://homepages.inf.ed.ac.uk/wadler/papers/expression/expression.txt> Discussion on the Java-genericity mailing list.
- [45] Yizhou Zhang and Andrew C. Myers. 2017. Familia: Unifying interfaces, type classes, and family polymorphism. *Proc. of the ACM on Programming Languages (PACMPL)* 1, OOPSLA (Oct. 2017). <https://doi.org/10.1145/3133894>